



STAT 311: LECTURE 13

Heavily based on lecture notes from Martina Morris



Random Variables

Logistics



- Homework due today
- Midterms to be passed back Friday
- Shiqing will be giving lectures on Monday and Wednesday
- Sam's office hours will be on Friday from 1-3



Random Variables

From probabilities of specific events
To describing the full probability distribution

What is a random variable?

Formally:

- An event in a sample space
- That takes a value
 - Discrete or
 - Continuous
- With a certain probability
 - Can be described by the PDF: $P(X=k)$
 - Or by the CDF: $P(X \leq k)$, $P(X > k)$, $P(j < X < k)$



Examples

- Flip a coin: Heads or tails?
- How many cars will drive through an intersection in an hour?
- What is the height of a random individual?
- How much time until I get my next text message?

Random variable notation

- Random variables are denoted by capital letters
 - For example, X or Y
- The value the RV takes in a specific case is called a “realization” and denoted by a lowercase letter
 - For example: x , y or k
- $P(X = k)$
 - “The probability that the random variable X takes the value k ”
- The sample space \mathcal{X} (set of all possible outcomes) is denoted by

Empirical vs Theoretical Distributions

- We have seen distributions of observed data. These are often called **Empirical Distributions**, because they are what we empirically observe
- Today we will begin to formally discuss distributions which are mathematical constructs used to model real world situations
- These distributions are typically a family of distributions which are specified by a mathematical equation and are governed by a set of **parameters**

Notation: statistics vs. parameters

Sample Statistics

- Mean

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$$

- Variance

$$s_x^2 = \frac{1}{(n-1)} \sum_{i=1}^n (x_i - \bar{x})^2$$

- Each observation has equal weight (1/n)

Theoretical parameters

- Expected value

$$\mu_x = \sum_{x_i \in \mathcal{X}} p_i x_i$$

- Variance

$$\sigma_x^2 = \sum_{x_i \in \mathcal{X}} p_i (x_i - \mu_x)^2$$

- Each possible outcome in the sample space receives its own weight (p_i)

Example

- If I roll a single fair dice

$$\mathcal{X} = \{1, 2, 3, 4, 5, 6\}$$

$$\begin{aligned}\mu_x &= \sum_{x_i \in \mathcal{X}} x_i p_i \\ &= 1(1/6) + 2(1/6) + 3(1/6) + 4(1/6) + 5(1/6) + 6(1/6) \\ &= 21/6 = 3.5\end{aligned}$$

$$\begin{aligned}\sigma_x^2 &= \sum_{x_i \in \mathcal{X}} (x_i - \mu_x)^2 p_i \\ &= (1 - 3.5)^2(1/6) + (2 - 3.5)^2(1/6) \dots (6 - 3.5)^2(1/6) = 2.92\end{aligned}$$

What comes next

- Deriving expectations and variances for different distributions
 - Discrete (Bernoulli, Binomial, Poisson)
 - Continuous (Normal, Uniform, Exponential)
- For example: with coin tosses
 - Each toss is a random variable with outcomes $\{0, 1\}$
 - The sum of these outcomes over n tosses is *a linear combination of the random variables* for each toss, with outcome space $\{0, 1, 2, \dots, n\}$
- We start with the rules of expectations and variances for linear transformations and combinations of RVs

Rules of expectations and variances

For linear transformations and combinations of random variables

1. Variance-Mean relationship

$$\begin{aligned}\sigma_x^2 = \text{var}(x) &= \sum_i p_i (x_i - \mu_x)^2 = \sum_{x_i \in \mathcal{X}} (x_i^2 - 2x_i\mu_x + \mu_x^2)p_i \\ &= \sum_{x_i \in \mathcal{X}} (x_i^2 p_i - 2x_i\mu_x p_i + \mu_x^2 p_i) \\ &= \sum_{x_i \in \mathcal{X}} x_i^2 p_i - \sum_{x_i \in \mathcal{X}} 2x_i\mu_x p_i + \sum_{x_i \in \mathcal{X}} \mu_x^2 p_i \\ &= \sum_{x_i \in \mathcal{X}} x_i^2 p_i - \mu_x \sum_{x_i \in \mathcal{X}} 2x_i p_i + \mu_x^2 \sum_{x_i \in \mathcal{X}} p_i \\ &= E(X^2) - 2(E(X))^2 + (E(X))^2 \\ &= E(X^2) - (E(X))^2\end{aligned}$$

Note the theoretical or population variance formula here, not the sample variance.

Transformations and combinations of RVs

■ Examples:

- *Transformation*: converting degrees F to degrees C
- *Combination*: adding your midterm and final exam scores

■ Linear transformations and combinations have simple expressions for their expected values and variances

- Linear transformations : $Y = a + X$, $Y = bX$, $Y = a + bX$,
- Linear combinations: $Z = X + Y$, $Z = aX + bY$

2. Linear Transformations of RVs

- If a and b are constants, and X is a random variable, then:

$$E(a) = a$$

$$\text{Var}(a) = 0$$

$$E(bX) = b E(X)$$

$$\text{Var}(bX) = b^2 \text{Var}(X)$$

$$E(a + bX) = a + b E(X)$$

$$\text{Var}(a + bX) = b^2 \text{Var}(X)$$

- These are easily proven if you start with the definitions on slide 8 and work through the algebra (*try it...*).

3. Linear combinations of independent RVs

UH covers this:

Let a , b and c be constants, and X and Y be independent random variables

$$E(X + Y) = E(X) + E(Y)$$

$$E(X - Y) = E(X) - E(Y)$$

$$E(a + bX + cY) = a + bE(X) + cE(Y)$$

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y)$$

$$\text{Var}(X - Y) = \text{Var}(X) + \text{Var}(Y)$$

$$\text{Var}(a + bX + cY) = b^2\text{Var}(X) + c^2\text{Var}(Y)$$

Again, these are easily proven

4. Linear combinations of dependent RVs

- Not covered in UH, but straightforward

Let a and b be constants, and X and Y be random variables

$$E(X + Y) = E(X) + E(Y)$$

No difference in the mean

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y) + 2\text{Cov}(X, Y)$$

$$\text{Var}(X - Y) = \text{Var}(X) + \text{Var}(Y) - 2\text{Cov}(X, Y)$$

$$\text{Var}(aX + bY) = a^2\text{Var}(X) + b^2\text{Var}(Y) + 2ab\text{Cov}(X, Y)$$

But the variance changes

Why the difference for correlated RVs?

- Suppose 10 individuals are deciding whether or not to show up to a party. Each has a 50/50 chance of going
- If the individuals all decide independently, we tend to end up with very few extreme events (ie 0 show up or all 10 show up) and usually will have around 5 individuals
- If the individuals all text each other and decide to either all show up or all not show up (attendance for each individual is now dependent), we will always have either 0 or 10 individuals. The average is still 5, but the outcomes are more extreme

Summary



- There are simple rules for expectations and variances
- When we transform or combine random variables
 - As long as the transformation/combination is linear
- And these are the foundation for what comes next
 - Derive expected values and variances for some common distributions

Discrete random variables

Exploring the derivation of discrete probability distributions and their properties:

Bernoulli, Binomial and Poisson

The goal

- To define a theoretical **Probability Density Function** for a random variable

“Probability that the RV $X=k$, given p ”

“The RV X is distributed as f with parameter p ”

Where p is one or more *parameters* that determine the outcome of the random variable.

- And use the PDF to derive expected values, variances, and probabilities

With discrete distributions, $f(x)$ is technically called a “*probability mass function*” or *PMF*, but PDF is also used, and we will use it here.

Example

- Coin tosses:
 - Each individual toss is an RV with 2 outcomes
 - Let X be the random variable for each toss: $\mathcal{X} = \{H, T\}$
- Pass Stat 311:
 - Each individual student is an RV with 2 outcomes
 - Let X be the random variable for each toss:
 $\mathcal{X} = \{\text{Pass}, \text{No Pass}\}$

The Bernoulli distribution

- Let X be a random variable with two outcomes:
 $\mathcal{X} = \{0, 1\}$ (we have to decide what is 1 and what is 0)

$$X \sim \text{Bernoulli}(p)$$

$$f(x; p) = p^x (1 - p)^{(1-x)}$$

$$f(1; p) = p^1 (1 - p)^{(1-1)} = p$$

$$f(0; p) = p^0 (1 - p)^{(1-0)} = 1 - p$$

Derivation of $E(Y)$ and $Var(Y)$

$X \sim \text{Bernoulli}(p)$

$$\begin{aligned}\mu_x &= \sum_{x_i \in \mathcal{X}} x_i p_i \\ &= 1(p) + 0(1-p) = p\end{aligned}$$

$$\begin{aligned}\sigma_x^2 &= \sum_{x_i \in \mathcal{X}} (x_i - \mu_x)^2 p_i \\ &= (1-p)^2(p) + (0-p)^2(1-p) \\ &= (1-2p+p^2)(p) + p^2(1-p) \\ &= (p-2p^2+p^3) + p^2 - p^3 \\ &= p - p^2 = p(1-p)\end{aligned}$$

Getting more complicated

- If there are 50 students in a class, how many total students will show up to class.
 - Assume each student shows up with probability .8
 - Assume the attendance of each student is independent of other students
- Each student's attendance is a Bernoulli trial
- Want to know the sum of the Bernoulli trials
- How can we do this?

Repeated Bernoulli trials: *the Binomial*

- Define a general probability distribution for the sum of n **independent** Bernoulli trials
- Let $X = \{0, 1, 2, \dots, n\}$ the count of the number of 1's.

$$X \sim \text{Binomial}(n; p)$$

Example: 3 Coin tosses with H=1

		<u>TTT</u>	<u>TTH</u>	<u>THT</u> HTT THT	<u>THH</u> HHT HTH	<u>HHH</u>
Value of X	0	1	2	3		
Probability	1/8	3/8	3/8	1/8		

The Binomial distribution

-
- **The Binomial describes the probability distribution of counts of successes/failures**
- It is a linear combination (sum) of Bernoulli RVs
- So both the number of trials, n , and the probability of each trial, p , outcome influence the result
- And the outcome space is now $\{0, 1, 2, \dots, n\}$

Binomial probabilities

There are 3 elements to the calculation:

1. Define the probability of each individual outcome in the sample space.
2. Identify the number of outcomes in the set of interest (i.e., that satisfy the condition $X=k$).
3. Multiply the probability of the outcome by the number in the set

What is the probability of each outcome?

- Start with a single trial:
 - What is the probability of each outcome Y?

$$p^y (1 - p)^{1-y}$$

Our friend the Bernoulli

- What about
 - What is the probability of each outcome

$$p^k (1 - p)^{2-k}$$

- And in general: $\dots p^k (1 - p)^{n-k}$

How many outcomes satisfy $X = k$

- This is a counting problem
- How many ways to get k successes in n trials when order does not matter?
- Out of our N trials, we need to choose k trials to be the successes
- Use the combination rule: $C_N^k = \binom{N}{k}$
- Use the combination rule:

The “binomial coefficient”

The number of outcomes that satisfy the condition:

$$\begin{array}{cccc} & & \text{HTT HHT} & \\ & & \text{THT HTH} & \\ & \text{TTT} & \text{TTH} & \text{THH HHH} \\ \binom{n}{k} = & \binom{3}{0}, & \binom{3}{1}, & \binom{3}{2}, & \binom{3}{3} \\ & = 1, & 3, & 3, & 1 \end{array}$$

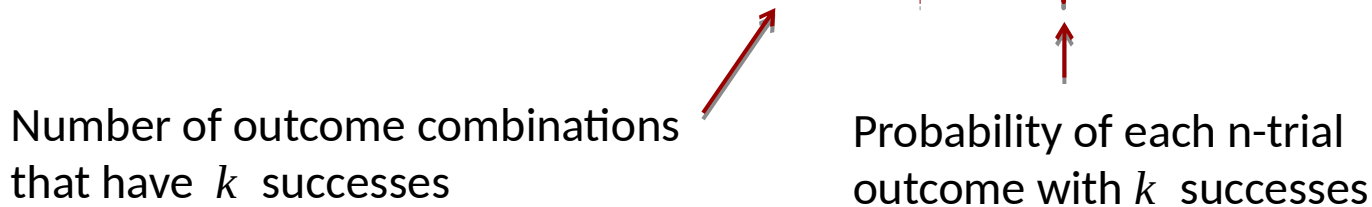
- is referred to as the binomial coefficient in this context

The Binomial distribution

Putting these all together

Let Y be a random variable with two outcomes, $Y=\{0,1\}$

Let X be the number of successes in n trials, and p be the probability of success on each trial. Then:

$$X \sim \text{Bin}(n; p) \quad P(X = k) = \binom{n}{k} \underbrace{p^k (1-p)^{n-k}}$$


Number of outcome combinations
that have k successes

Probability of each n -trial
outcome with k successes

Derivation of $E(X)$ and $Var(X)$

$X \sim \text{Bin}(n; p)$ the sum of n independent Bernoulli trials: $X = \sum_{i=1}^n Y_i$

$$E(X) = \mu_X$$

$$= E\left(\sum_{i=1}^n Y_i\right)$$

$$= \sum_{i=1}^n E(Y_i) *$$

$$= n\mu_Y$$

$$= np$$

$$Var(X) = \sigma_X^2$$

$$= Var\left(\sum_{i=1}^n Y_i\right)$$

$$= n\sigma_Y^2 *$$

$$= np(1-p)$$

* Using $E(Z+W) = E(Z) + E(W)$

* Using $Var(Z+W) = Var(Z) + Var(W)$
for independent RVs Z and W

Binomial Summary

- For **any** repeated Bernoulli trial, the count of successes, X , has a Binomial distribution:

$$X \sim \text{Bin}(n; p)$$

$$f(x; n, p) = \binom{n}{x} p^x (1-p)^{n-x}$$

$$\mu_X = np$$

$$\sigma_X^2 = np(1-p)$$

We can calculate the mean, variance, and the probability of any value of X from just two values: n and p

Other models

- Suppose I know that on average 10 cars pass through a specific intersection near my house each hour
- The number of cars that pass through for a given hour is random
- It is discrete (we're only considering whole cars)
- No set number of trials, or maximum value that this rv can take
- How might we describe this process?

Poisson Distribution


■ Poisson – Used for counts of events as rates

- n (the number of trials) is large, not fixed,
- λ approximates a rate of events (e.g., per time unit, or per capita)

for $x \geq 0$

$$f(x; \lambda) = P(X = x) = \frac{e^{-\lambda} \lambda^x}{x!}$$

$$\mu_X = \sigma_X^2 = \lambda$$



Poisson Distribution

- The numerator is always positive, so there is a positive probability for all $x \geq 0$, although it gets very small as x becomes large
- We assume that there is a constant rate
- We assume that all “arrivals” are independent of each other

Example

- Suppose I know that on average 10 cars pass through a specific intersection near my house each day

- What is λ ?
- What is the probability that 8 cars pass through the intersection in an hour?

$$f(8; \lambda = 10) = \frac{\lambda^k e^{-\lambda}}{k!} = \frac{10^8 e^{-10}}{8!} = .1125$$

- What is the probability that 15 cars pass through the intersection in an hour?

$$f(15; \lambda = 10) = \frac{\lambda^k e^{-\lambda}}{k!} = \frac{10^{15} e^{-10}}{15!} = .0347$$

- Why might a Poisson assumption be wrong for this model?

Summary:

- A discrete random variable takes specific values
 - Typically integer counts
 - Each with a certain probability
- If you can represent the underlying stochastic process as a mathematical function
 - You can calculate almost anything you want for the RV
 - $E(X)$, $\text{Var}(X)$, $P(X=k)$, $P(X \leq k)$ or $P(i < X < j)$
- The distribution is defined by the stochastic process
 - The details of the process matter, and are reflected in the formal definition of the distribution